

# ELEN\_6885\_HW4\_Part\_1\_2\_3

November 19, 2019

## 1 ELEN 6885 Reinforcement Learning Coding Assignment (Part 1, 2, 3)

### 1.1 Taxi Problem Overview

Please put your code into the block marked by: #####  
YOUR CODE STARTS HERE YOUR CODE ENDS HERE  
##### You should not edit anything outside  
of the block.

[0]:

## 2 Playing with the environment

Run the cell below to get a feel for the environment by moving your agent(the taxi) by taking one of the actions at each step.

```
[0]: from gym.wrappers import Monitor
import gym
import random
import numpy as np
```

```
[0]: """
You can test your game now.
Input range from 0 to 5:
    0 : South (Down)
    1 : North (Up)
    2 : East (Right)
    3 : West (Left)
    4: Pick up
    5: Drop off
    6: exit_game
"""
GAME = "Taxi-v3"
env = gym.make(GAME)
env = Monitor(env, "taxi_simple", force=True)
s = env.reset()
steps = 100
```

```

for step in range(steps):
    env.render()
    action = int(input("Please type in the next action:"))
    if action==6:
        break
    s, r, done, info = env.step(action)
    print('state:',s)
    print('reward:',r)
    print('Is state terminal?:',done)
    print('info:',info)

# close environment and monitor
env.close()

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

Please type in the next action:6

## 2.1 1.1 Incremental implementation of average

We've finished the incremental implementation of average for you. Please call the function to estimate with  $1/\text{step}$  step size and fixed step size to compare the difference between these two on a simulated Bandit problem.

```

[0]: def estimate(OldEstimate, StepSize, Target):
    '''An incremental implementation of average.
    OldEstimate : float
    StepSize : float
    Target : float
    '''
    NewEstimate = OldEstimate + StepSize * (Target - OldEstimate)
    return NewEstimate

```

```

[0]: random.seed(6885)
numTimeStep = 10000
q_h = np.zeros(numTimeStep + 1) # Q Value estimate with 1/step step size
q_f = np.zeros(numTimeStep + 1) # Q value estimate with fixed step size
FixedStepSize = 0.5 #A large number to exaggerate the difference
for step in range(1, numTimeStep + 1):
    if step < numTimeStep / 2:
        r = random.gauss(mu = 1, sigma = 0.1)
    else:

```

```

r = random.gauss(mu = 3, sigma = 0.1)

#TIPS: Call function estimate defined in ./helpers/utils.py
#####
# YOUR CODE STARTS HERE
q_f[step]=estimate(q_f[step-1],FixedStepSize,r)
q_h[step]=estimate(q_h[step-1],1/step,r)

# YOUR CODE ENDS HERE
#####

q_h = q_h[1:]
q_f = q_f[1:]

```

Plot the two Q value estimates. (Please include a title, labels on both axes, and legends)

```

[0]: import matplotlib.pyplot as plt
#####
# YOUR CODE STARTS HERE
plt.plot(q_h)
plt.plot(q_f)
plt.xlabel('step')
plt.ylabel('q-value estimate')
plt.title('fixed stepsize vs 1/step stepsize')
plt.legend(['1/step', 'fixed'])

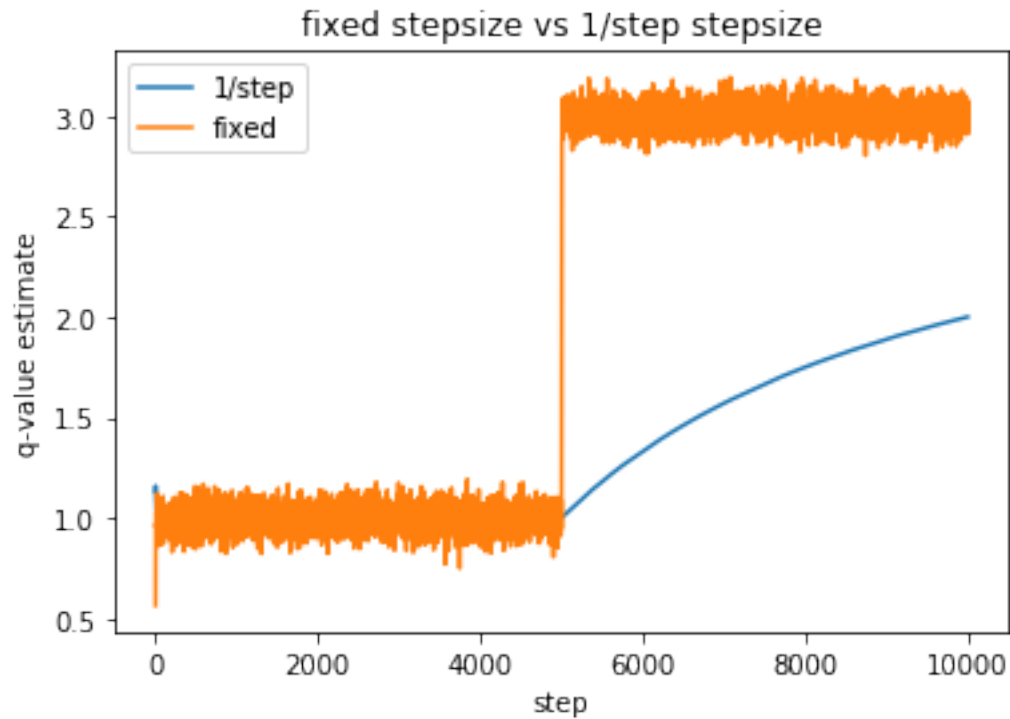
# YOUR CODE ENDS HERE
#####

```

```

[0]: <matplotlib.legend.Legend at 0x7ff13326a908>

```



## 2.2 1.2 $\epsilon$ -Greedy for Exploration

In Reinforcement Learning, we are always faced with the dilemma of exploration and exploitation.  $\epsilon$ -Greedy is a trade-off between them. You are supposed to implement Greedy and  $\epsilon$ -Greedy. We combine these two policies in one function by treating Greedy as  $\epsilon$ -Greedy where  $\epsilon = 0$ . Edit the function `epsilon_greedy` the following block.

```
[0]: def epsilon_greedy(value, e, seed = None):
    """
    Implement Epsilon-Greedy policy.

    Inputs:
    value: numpy ndarray
    A vector of values of actions to choose from
    e: float
    Epsilon
    seed: None or int
    Assign an integer value to remove the randomness

    Outputs:
    action: int
    Index of the chosen action
    """
    assert len(value.shape) == 1
```

```

assert 0 <= e <= 1

if seed != None:
    np.random.seed(seed)

#####
# YOUR CODE STARTS HERE
n=len(value)
#####
# YOUR CODE STARTS HERE
np.random.seed(0)
randProb = np.random.random() # Pick random probability between 0-1
if randProb < e:
    a = np.random.choice(len(value)) # Select random action
else:
    maxAction = np.argmax(value) # Find max value estimate
    action = np.where(value == np.argmax(value))[0]
    if len(action) == 0:
        a = maxAction
    else:
        a = np.random.choice(action)

# YOUR CODE ENDS HERE
#####
return a

```

```

[0]: np.random.seed(6885) #Set the seed for reproducibility
q = np.random.normal(0, 1, size = 5)
#####
# YOUR CODE STARTS HERE
greedy_action=epsilon_greedy(q, 0, seed = None)
e_greedy_action=epsilon_greedy(q, 0, seed = None)
# YOUR CODE ENDS HERE
#####
print('Values:')
print(q)
print('Greedy Choice =', greedy_action)
print('Epsilon-Greedy Choice =', e_greedy_action)

```

Values:

[ 0.61264537 0.27923079 -0.84600857 0.05469574 -1.09990968]

Greedy Choice = 0

Epsilon-Greedy Choice = 0

You should get the following results: Values: [ 0.61264537 0.27923079 -0.84600857 0.05469574 -1.09990968] Greedy Choice = 0 Epsilon-Greedy Choice = 0

## 2.3 1.3 Exploration VS. Exploitation

Try to reproduce Figure 2.2 (the upper one is enough) of the Sutton's book based on the experiment described in Chapter 2.3.

```
[0]: # Do the experiment and record average reward acquired in each time step
#####
# YOUR CODE STARTS HERE
import numpy as np
import matplotlib.pyplot as plt

Arms = 10 # n number of bandits
iterations = 2000 # number of repeated iterations
plays = 1000
kAction = np.zeros(Arms) # count of actions taken at time t
rSum = np.zeros(Arms)

valEstimates = np.zeros(Arms)

scoreArr = np.zeros(plays)
scoreArr1 = np.zeros(plays)
scoreArr2 = np.zeros(plays)
rAvg=np.zeros(plays)

def action(valEstimates,e):
    randProb = np.random.random() # Pick random probability between 0-1
    if randProb < e:
        a = np.random.choice(len(valEstimates)) # Select random action
    else:
        maxAction = np.argmax(valEstimates) # Find max value estimate
        action = np.where(valEstimates == np.argmax(valEstimates))[0]
        if len(action) == 0:
            a = maxAction
        else:
            a = np.random.choice(action)
    return a

for iIter in range(iterations):

    q_star = np.random.normal(0, 1, 10)
    #actionT = None # Store last action
    kAction = np.zeros(Arms)
    kAction1 = np.zeros(Arms)
    kAction2 = np.zeros(Arms) # count of actions taken at time t
    rSum = np.zeros(Arms)
    rSum1 = np.zeros(Arms)
    rSum2 = np.zeros(Arms)
```

```

valEstimates = np.zeros(Arms)
valEstimates1 = np.zeros(Arms)
valEstimates2 = np.zeros(Arms)
if (iIter % 100) == 0:
    print("Completed Iterations: ", iIter)
for jPlays in range(plays):
    actionT = action(valEstimates,e=0)
    actionT1 = action(valEstimates1, e=0.1)
    actionT2 = action(valEstimates2, e=0.01)
    rewardT = np.random.normal(q_star[actionT], 1)
    rewardT1 = np.random.normal(q_star[actionT1], 1)
    rewardT2 = np.random.normal(q_star[actionT2], 1)
    At = actionT
    At1 = actionT1
    At2 = actionT2
    kAction[At] += 1
    kAction1[At1] += 1
    kAction2[At2] += 1# Add 1 to action selection
    rSum[At] += rewardT
    rSum1[At1] += rewardT1
    rSum2[At2] += rewardT2# Add reward to sum array
    valEstimates[At] = rSum[At] / kAction[At]
    valEstimates1[At1] = rSum1[At1] / kAction1[At1]
    valEstimates2[At2] = rSum2[At2] / kAction2[At2]
    scoreArr[jPlays] += rewardT
    scoreArr1[jPlays] += rewardT1
    scoreArr2[jPlays] += rewardT2
rAvg = scoreArr / iterations
rAvg1 = scoreArr1/ iterations
rAvg2 = scoreArr2 / iterations

# YOUR CODE ENDS HERE
#####

```

```

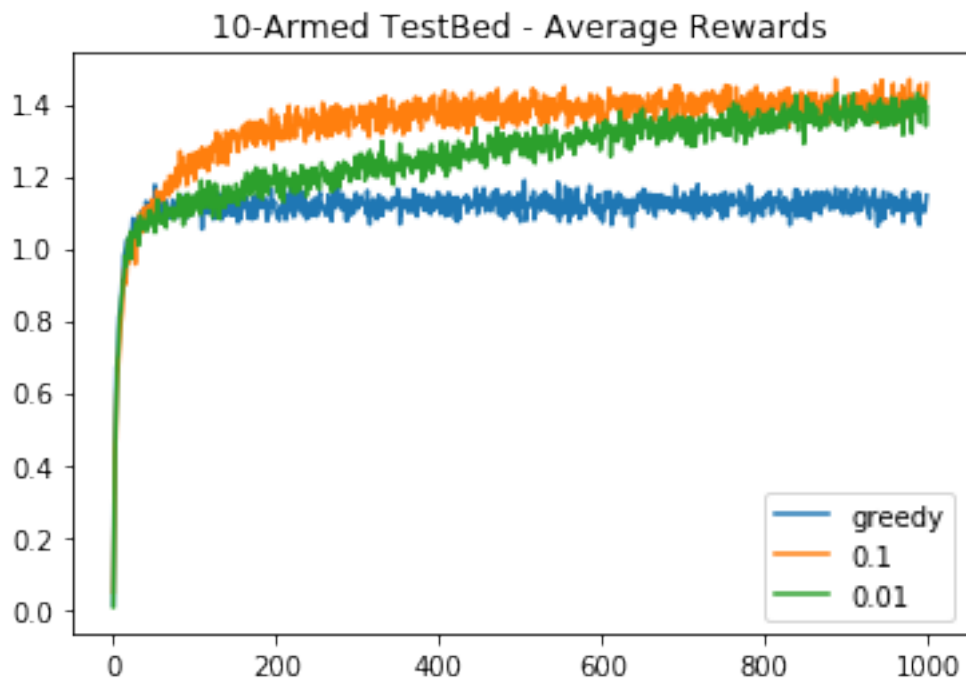
Completed Iterations: 0
Completed Iterations: 100
Completed Iterations: 200
Completed Iterations: 300
Completed Iterations: 400
Completed Iterations: 500
Completed Iterations: 600
Completed Iterations: 700
Completed Iterations: 800
Completed Iterations: 900
Completed Iterations: 1000
Completed Iterations: 1100
Completed Iterations: 1200

```

Completed Iterations: 1300  
Completed Iterations: 1400  
Completed Iterations: 1500  
Completed Iterations: 1600  
Completed Iterations: 1700  
Completed Iterations: 1800  
Completed Iterations: 1900

```
[0]: # Plot the average reward
#####
# YOUR CODE STARTS HERE
plt.title("10-Armed TestBed - Average Rewards")
plt.plot(rAvg)
plt.plot(rAvg1)
plt.plot(rAvg2)
plt.legend(['greedy', '0.1', '0.01'], loc=4)
plt.show()

# YOUR CODE ENDS HERE
#####
```





## 3 Question 2

In this question, you will implement the value iteration and policy iteration algorithms to solve the Taxi game problem

### 3.1 2.1 Model-based RL: value iteration

For this part, you need to implement the helper functions `action_evaluation(env, gamma, v)`, and `extract_policy(env, v, gamma)` in `utils.py`. Understand `action_selection(q)` which we have implemented. Use these helper functions to implement the `value_iteration` algorithm below.

```
[0]: import numpy as np
from helpers import utils
def value_iteration(env, gamma, max_iteration, theta):
    """
    Implement value iteration algorithm. You should use extract_policy to for
    ↪ extracting the policy.

    Parameters
    -----
    env: OpenAI env.
        env.P: dictionary
            the transition probabilities of the environment
            P[state][action] is tuples with (probability, nextstate,
    ↪ reward, terminal)
        env.nS: int
            number of states
        env.nA: int
            number of actions

    gamma: float
        Discount factor.
    max_iteration: int
        The maximum number of iterations to run before stopping.
    theta: float
        Determines when value function has converged.

    Returns:
    -----
    value function: np.ndarray
    policy: np.ndarray
    """
    nS = env.nS
    nA = env.nA
    V = np.zeros(env.nS)
    #####
    # YOUR CODE STARTS HERE
    for i in range(max_iteration):
        q=utils.action_evaluation(env, gamma, V)
        #print(q)
```

```

policy=np.argmax(q, axis = 1)
#q=action_evaluation(env, gamma, v)
delta=[]
for s in range(nS):
    temp=V[s]
    #print(temp)
    V[s]=max(q[s][:])
    #print(V[s])
    #print(abs(temp-V[s]))
    delta.append(abs(temp-V[s]))
if max(delta)<theta:
    break
# YOUR CODE ENDS HERE
#####

return V, policy

```

After implementing the above function, read and understand the functions implemented in evaluation\_utils.py, which we will use to evaluate our value iteration policy

```

[0]: from helpers import evaluation_utils
import gym
GAME = "Taxi-v3"
env = gym.make(GAME)
V_vi, policy_vi = value_iteration(env, gamma=0.95, max_iteration=6000,
    ↪theta=1e-5)
# visualize how the agent performs with the policy generated from value_
    ↪iteration
evaluation_utils.render_episode(env, policy_vi)

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(South)
+-----+

```

```
|R: | : :G| |
| : | : : |
| : : | : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(West)

```
+-----+
```

```
|R: | : :G| |
| : | : : |
| : | : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(West)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(West)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| : | : | : |
|Y| : |B: |
```

```
+-----+
```

(South)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(South)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(Pickup)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : : : : |
|█| : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : :G|
| : | : : |
|█: : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| :█: : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :G|
| : | : : |
| : :█: : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :G|
| : |█: : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: |█: :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(North)

```
+-----+
```

```
|R: | : █:G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :█|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(East)

```
+-----+
```

```
|R: | : :█|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
```

```
+-----+
```

(Dropoff)

Episode reward: 5.000000

```
[0]: # evaluate the performance of value iteration over 100 episodes
evaluation_utils.avg_performance(env, policy_vi)
```

```
[0]: 8.121212121212121
```

## 3.2 2.2 Model-based RL: policy iteration

In this part, you are supposed to implement policy iteration to solve the Taxi game problem.

```
[0]: #from helpers import utils
def policy_iteration(env, gamma, max_iteration, theta):
    """Implement Policy iteration algorithm.

    You should use the policy_evaluation and policy_improvement methods to
    implement this method.

    Parameters
    -----
    env: OpenAI env.
        env.P: dictionary
            the transition probabilities of the environment
            P[state][action] is tuples with (probability, nextstate,
            ↪reward, terminal)
```

```

        env.nS: int
            number of states
        env.nA: int
            number of actions
    gamma: float
        Discount factor.
    max_iteration: int
        The maximum number of iterations to run before stopping.
    theta: float
        Determines when value function has converged.
    Returns:
    -----
    value function: np.ndarray
    policy: np.ndarray
    """

    V0= np.zeros(env.nS)
    policy = np.zeros(env.nS, dtype=int)
    #####
    # YOUR CODE STARTS HERE
    nS = env.nS
    nA = env.nA
    policy_stable=['b']
    while True:
        if 'b' not in policy_stable:
            break
        else:
            #delta=[]
            #for i in range(max_iteration):
            #for s in range(nS):
            #temp1=V[s]
            #V = policy_evaluation(env, policy, gamma, theta)
            #delta.append(abs(temp1 - V[s]))
            #if max(delta) < theta:
            #break
            V = policy_evaluation(env, policy, gamma, theta,max_iteration,V0)
            policy, policy_stable = policy_improvement(env, V, policy, gamma)

    # YOUR CODE ENDS HERE
    #####

    return V, policy

def policy_evaluation(env, policy, gamma, theta,max_iteration,V):
    """Evaluate the value function from a given policy.

```

```

Parameters
-----
env: OpenAI env.
    env.P: dictionary
        the transition probabilities of the environment
        P[state][action] is tuples with (probability, nextstate,
→reward, terminal)
    env.nS: int
        number of states
    env.nA: int
        number of actions

gamma: float
    Discount factor.
policy: np.array
    The policy to evaluate. Maps states to actions.
max_iteration: int
    The maximum number of iterations to run before stopping.
theta: float
    Determines when value function has converged.
Returns
-----
value function: np.ndarray
    The value function from the given policy.
"""
#V = np.zeros(env.nS)###
#####
# YOUR CODE STARTS HERE
nS = env.nS
nA = env.nA
#v = np.zeros(nS)
P = env.P
for i in range(max_iteration):
    delta = []
    for s in range(nS):
        temp1 = V[s]
        v_s = 0
        for i in range(len(P[s][policy[s]])):
            next_state_tuple = P[s][policy[s]][i]
            v_next_state = V[next_state_tuple[1]]
            p_next_state = next_state_tuple[0]
            reward_next_state = next_state_tuple[2]
            v_s += p_next_state * (reward_next_state + gamma * v_next_state)
            delta.append(abs(temp1 - V[s]))
        V[s] = v_s
    if max(delta) < theta:

```

```

        break

    # YOUR CODE ENDS HERE
    #####

    return V

def policy_improvement(env, value_from_policy, policy, gamma):
    """Given the value function from policy, improve the policy.

    Parameters
    -----
    env: OpenAI env
        env.P: dictionary
            the transition probabilities of the environment
            P[state][action] is tuples with (probability, nextstate, ↪reward, terminal)
        env.nS: int
            number of states
        env.nA: int
            number of actions

    value_from_policy: np.ndarray
        The value calculated from the policy
    policy: np.array
        The previous policy.
    gamma: float
        Discount factor.

    Returns
    -----
    new_policy: np.ndarray
        An array of integers. Each integer is the optimal action to take
        in that state according to the environment dynamics and the
        given value function.
    stable_policy: bool
        True if the optimal policy is found, otherwise false
    """
    #####
    # YOUR CODE STARTS HERE
    nS = env.nS
    nA = env.nA
    q = np.zeros((nS, nA))
    new_policy = np.zeros(env.nS, dtype=int)
    P = env.P
    temp=[]

```



```

policy_stable = []
for s in range(nS):
    for a in range(nA):
        q_s_a = 0
        for i in range(len(P[s][a])):
            next_state_tuple = P[s][a][i]
            v_next_state = value_from_policy[next_state_tuple[1]]
            p_next_state = next_state_tuple[0]
            reward_next_state = next_state_tuple[2]
            q_s_a += p_next_state * (reward_next_state + gamma *
↪v_next_state)
        q[s][a] = q_s_a
        temp.append(policy[s])
        new_policy[s] = int(np.argmax(q[s]))
        if new_policy[s] == temp[s]:
            policy_stable.append('a')
        else:
            policy_stable.append('b')
# print(new_policy)
# YOUR CODE ENDS HERE
#####

return new_policy, policy_stable

```

```

[0]: ## Testing out policy iteration policy for one episode
from helpers import evaluation_utils
import gym
GAME = "Taxi-v3"
env = gym.make(GAME)
#evaluation_utils.render_episode(env, policy_vi)
V_pi, policy_pi = policy_iteration(env, gamma=0.95, max_iteration=6000,
↪theta=1e-5)

```

```

[0]: # visualize how the agent performs with the policy generated from policy
↪iteration
evaluation_utils.render_episode(env, policy_pi)

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

```

+-----+
|R: | : :G|

```

```

| : | : : |
| : : : : |
| | : | : |
|Y| : | |B: |

```

```

+-----+

```

(East)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | | : |
|Y| : |B: |

```

```

+-----+

```

(North)

```

+-----+

```

```

|R: | : :G| |
| : | : : |
| : : | : : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(North)

```

+-----+

```

```

|R: | : :G| |
| : | : : |
| : : : | : |
| | : | : |
|Y| : |B: |

```

```

+-----+

```

(East)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | | : |
|Y| : |B: |

```

```

+-----+

```

(South)

```

+-----+

```

```

|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : | B : |

```

```

+-----+

```

(South)

```

+-----+

```

```

|R: | : :G|

```

```

| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Pickup)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : |B: |
|Y| : |B: |
+-----+
(North)
+-----+
|R: | : :G|
| : | : : |
| : : B: |
| | : | : |
|Y| : |B: |
+-----+
(North)
+-----+
|R: | : :G|
| : | : : |
| : B: : : |
| | : | : |
|Y| : |B: |
+-----+
(West)
+-----+
|R: | : :G|
| : | : : |
| : B: : : |
| | : | : |
|Y| : |B: |
+-----+
(West)
+-----+
|R: | : :G|
| : B: | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(North)
+-----+
|R: B: | : :G|

```

```

| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
      (North)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
      (West)
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
      (Dropoff)

```

Episode reward: 6.000000

```
[0]: # evaluate the performance of policy iteration over 100 episodes
print(evaluation_utils.avg_performance(env, policy_pi))
```

8.383838383838384

## 4 Part 3: Q-learning and SARSA

### 4.1 3.1 Model-free RL: Q-learning

In this part, you will implement Q-learning.

```
[0]: def epsilon_greedy(value, e, seed=None):
    assert len(value.shape) == 1

    assert 0 <= e <= 1

    if seed != None:
        np.random.seed(seed)
    #n = len(value)
    #np.random.seed(0)
    randProb = np.random.random() # Pick random probability between 0-1
    if randProb < e:
        a = np.random.choice(len(value)) # Select random action
```

```

else:
    maxAction = np.argmax(value)  # Find max value estimate
    action = np.where(value == np.argmax(value))[0]
    if len(action) == 0:
        a = maxAction
    else:
        a = np.random.choice(action)
return a

def QLearning(env, num_episodes, gamma, lr, e):
    """
    Implement the Q-learning algorithm following the epsilon-greedy exploration.
    Inputs:
    env: OpenAI Gym environment
        env.P: dictionary
            P[state][action] are tuples of tuples tuples with
            ↪ (probability, nextstate, reward, terminal)
                probability: float
                nextstate: int
                reward: float
                terminal: boolean
        env.nS: int
            number of states
        env.nA: int
            number of actions
    num_episodes: int
        Number of episodes of training
    gamma: float
        Discount factor.
    lr: float
        Learning rate.
    e: float
        Epsilon value used in the epsilon-greedy method.
    Outputs:
    Q: numpy.ndarray
    """
    nS = env.nS
    P = env.P
    nA = env.nA
    Q = np.zeros((env.nS, env.nA))
    reward_com=[]
    #####
    # YOUR CODE STARTS HERE
    for i in range(num_episodes):
        #print(Q)

```

```

reward=0
state = np.random.randint(0, nS)
states=0
while True:

    action = epsilon_greedy(Q[state], e,seed=None)
    n0=len(P[state][action])
    p_next_state = np.zeros(len(P[state][action]),dtype=int)
    next_state=np.zeros(len(P[state][action]),dtype=int)
    for i in range(len(P[state][action])):
        next_state_tuple = P[state][action][i]
        next_state[i]=next_state_tuple[1]
        #print('$$',next_state_tuple[2])
        p_next_state[i] = next_state_tuple[0]
        #reward_next_state = next_state_tuple[2]

    new_state = np.random.choice(next_state, p=p_next_state.ravel())
    num=int(np.argmax(next_state==new_state))
    reward_next_state=P[state][action][num][2]
    states+=1
    terminal_state = P[state][action][num][3]
    Q[state][action] += lr * (reward_next_state + gamma *
    ↪max(Q[new_state])- Q[state][action])
    reward+=reward_next_state
    state = new_state
    if terminal_state==True:
        reward_ave=reward/states
        break

    reward_com.append(reward_ave)

# YOUR CODE ENDS HERE
#####
    #print(reward_com)
    return Q,reward_com

# YOUR CODE ENDS HERE
#####

```

[255]:

```

import matplotlib.pyplot as plt
Q,reward_com = QLearning(env = env.env, num_episodes = 1000, gamma = 1, lr = 0.
    ↪1, e = 0.1)

print('Action values:')
print(Q)

#####
# YOUR CODE STARTS HERE
plt.plot(reward_com)
plt.xlabel('episodes')
plt.ylabel('ave reward')
plt.title("Qlearning")

```

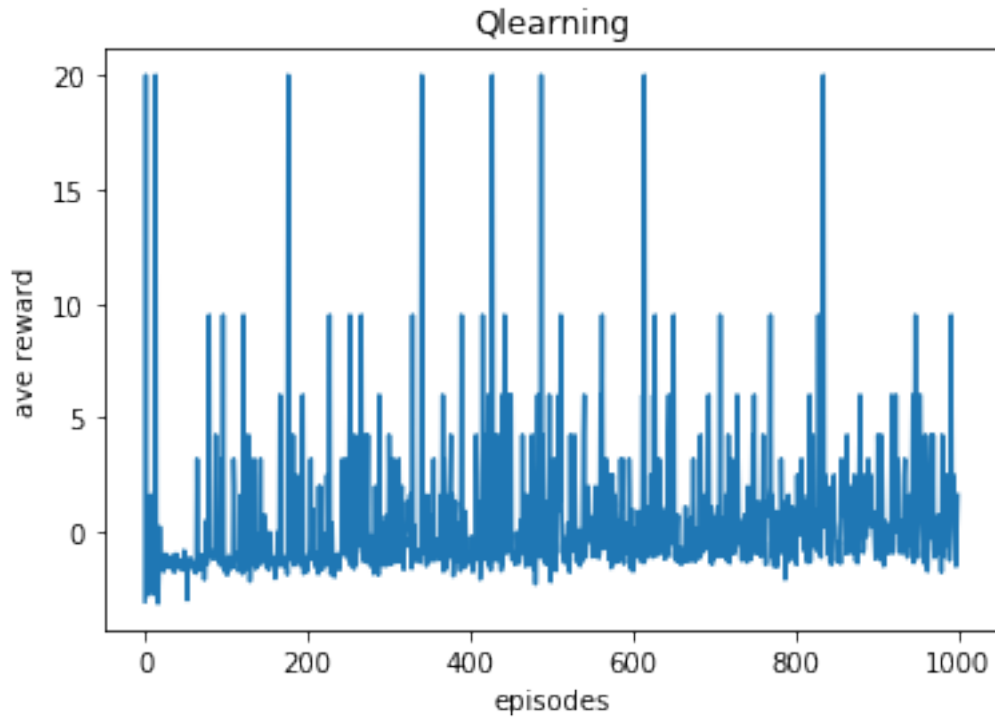
Action values:

```

[[ 5.19337550e+00 -1.00000000e-01 -1.00000000e-01  3.08443570e+00
   9.13028229e+01  0.00000000e+00]
 [-3.23969166e+00 -3.58930933e+00 -3.67201534e+00 -3.55985212e+00
   9.53536173e+00 -6.29542100e+00]
 [-1.03280583e+00 -3.29505878e-02 -1.45270334e+00 -1.57747241e-01
   3.14578010e+01 -3.47432407e-03]
 ...
 [-8.93662800e-01 -4.37554944e-01 -8.00000000e-01 -8.38818893e-01
  -1.00000000e+00 -1.91000000e+00]
 [-2.60000000e+00 -2.55861093e+00 -2.58903520e+00 -2.59902722e+00
  -3.00000000e+00 -3.00000000e+00]
 [ 0.00000000e+00 -1.10000000e-01 -1.00000000e-01  1.58556448e+01
   0.00000000e+00 -1.00000000e+00]]

```

[255]: Text(0.5, 1.0, 'Qlearning')



```
[2]: #Uncomment the following to evaluate your result, comment them when you
      ↳ generate the pdf
      #from helpers import utils

      #env = gym.make('Taxi-v3')
      #policy_estimate = utils.action_selection(Q)

      #evaluation_utils.render_episode(env, policy_estimate)
```

## 4.2 3.2 Model-free RL: SARSA

In this part, you will implement Sarsa.

```
[0]: def SARSA(env, num_episodes, gamma, lr, e):
      """
      Implement the SARSA algorithm following epsilon-greedy exploration.
      Inputs:
      env: OpenAI Gym environment
          env.P: dictionary
              P[state][action] are tuples of tuples tuples with
      ↳ (probability, nextstate, reward, terminal)
              probability: float
              nextstate: int
```



```

        reward: float
        terminal: boolean
    env.nS: int
        number of states
    env.nA: int
        number of actions
num_episodes: int
    Number of episodes of training
gamma: float
    Discount factor.
lr: float
    Learning rate.
e: float
    Epsilon value used in the epsilon-greedy method.
Outputs:
Q: numpy.ndarray
    State-action values
"""
nS = env.nS
P = env.P
nA = env.nA
Q = np.zeros((env.nS, env.nA))
#####
# YOUR CODE STARTS HERE
reward_com = []
for i in range(num_episodes):
    #print(Q)
    reward = 0
    states=0
    state = np.random.randint(0, nS)
    while True:
        action = epsilon_greedy(Q[state], e, seed=None)
        n0=len(P[state][action])
        p_next_state = np.zeros(len(P[state][action]), dtype=int)
        next_state=np.zeros(len(P[state][action]), dtype=int)
        for i in range(len(P[state][action])):
            next_state_tuple = P[state][action][i]
            next_state[i]=next_state_tuple[1]
            p_next_state[i] = next_state_tuple[0]
            #reward_next_state = next_state_tuple[2]

        new_state = np.random.choice(next_state, p=p_next_state.ravel())
        num=int(np.argmax(next_state==new_state))
        #print(state)
        reward_next_state=P[state][action][num][2]
        terminal_state = P[state][action][num][3]
        #print( reward_next_state)

```

```

        action_next=epsilon_greedy(Q[new_state], e,seed=None)
        Q[state][action] += lr * (reward_next_state + gamma *
↪Q[new_state][action_next]- Q[state][action])
        #Q[state][action] += lr * (reward_next_state + gamma *
↪max(Q[new_state])- Q[state][action])
        #print(Q[state][action])
        #print(Q[new_state][action_next])
        #pdb.set_trace()
        reward += reward_next_state
        states += 1
        state = new_state
        if terminal_state==True:
            reward_ave = reward / states
            break

        reward_com.append(reward_ave)
        # YOUR CODE ENDS HERE
        #####

    return Q,reward_com

def epsilon_greedy(value, e, seed=None):
    assert len(value.shape) == 1

    assert 0 <= e <= 1

    if seed != None:
        np.random.seed(seed)
    #n = len(value)
    #np.random.seed(0)
    randProb = np.random.random() # Pick random probability between 0-1
    if randProb < e:
        a = np.random.choice(len(value)) # Select random action
    else:
        maxAction = np.argmax(value) # Find max value estimate
        action = np.where(value == np.argmax(value))[0]
        if len(action) == 0:
            a = maxAction
        else:
            a = np.random.choice(action)
    return a

```

```

[0]: def render_episode_Q(env, Q):
    """Renders one episode for Q functionon environment.

    Parameters

```

```

-----
env: gym.core.Environment
    Environment to play Q function on.
Q: np.array of shape [env.nS x env.nA]
    state-action values.
"""

episode_reward = 0
state = env.reset()
done = False
while not done:
    env.render()
    time.sleep(0.5)
    action = np.argmax(Q[state])
    state, reward, done, _ = env.step(action)
    episode_reward += reward

print ("Episode reward: %f" %episode_reward)

```

```

[0]: Q, reward_com = SARSA(env = env.env, num_episodes = 1000, gamma = 1, lr = 0.1, e_
    ↪ = 0.1)
print('Action values:')
print(Q)
plt.plot(reward_com)
plt.xlabel('episodes')
plt.ylabel('ave reward')
plt.title("SARSA")

```

Action values:

```

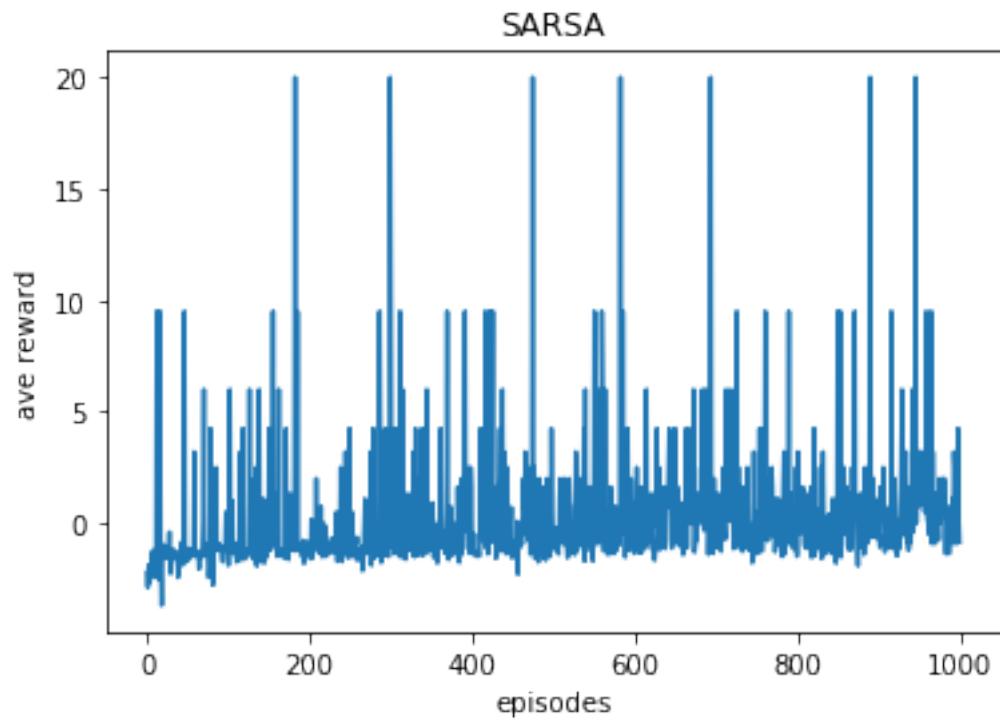
[[ 3.13426689e+00  2.78800349e+00  1.63642898e-02 -2.00000000e-01
   4.31518569e+01  1.80921795e+00]
 [-3.82198590e+00 -3.76951696e+00 -3.83169281e+00 -3.74401558e+00
  -1.12011272e+00 -4.79122233e+00]
 [-1.49408909e+00 -1.13406515e-01 -1.44254142e+00 -1.37799730e+00
   2.64419050e+01 -1.54021897e+00]
 ...
 [-1.00000000e+00 -9.12288460e-01 -9.90000000e-01  3.57393073e+00
  -1.98531261e+00 -1.00000000e+00]
 [-2.86697419e+00 -2.85807880e+00 -2.91009163e+00 -5.78758333e-01
  -3.97074853e+00 -3.00000000e+00]
 [-2.00000000e-01 -1.86254179e-01 -2.00000000e-01  1.78387360e+01
  -1.00000000e+00 -1.00000000e+00]]

```

```

[0]: Text(0.5, 1.0, 'SARSA')

```



```
[1]: # Uncomment the following to evaluate your result, comment them when you
      ↪ generate the pdf
      #from helpers import utils
      #env = gym.make('Taxi-v3')
      #policy_estimate = utils.action_selection(Q)
      #evaluation_utils.render_episode(env, policy_estimate)
```

```
[ ]:
```